



# An Introduction to Vitis High Level Synthesis

Duc Tri Nguyen



## Algorithm

$$s^2 = \sum_{i=0}^{n-1} x_i^2 = x_0^2 + x_1^2 + x_2^2 + \dots + x_{n-1}^2 \quad (1)$$

Equation 1 compute  $l^2$ -norm of a vector  $x = (x_0, x_1, \dots, x_{n-1})$ .

However, due to bit-width limitation, and the size of vector  $x$  is undefined, it can be arbitrary long or short. We use must use **saturation arithmetic** to perform add or subtract.



## Algorithm

$$s^2 = \sum_{i=0}^{n-1} x_i^2 = x_0^2 + x_1^2 + x_2^2 + \dots + x_{n-1}^2 \quad (1)$$

Can we optimize the algorithm?



# Algorithm

$$s^2 = \sum_{i=0}^{n-1} x_i^2 = x_0^2 + x_1^2 + x_2^2 + \dots + x_{n-1}^2 \quad (1)$$

Can we optimize the algorithm?

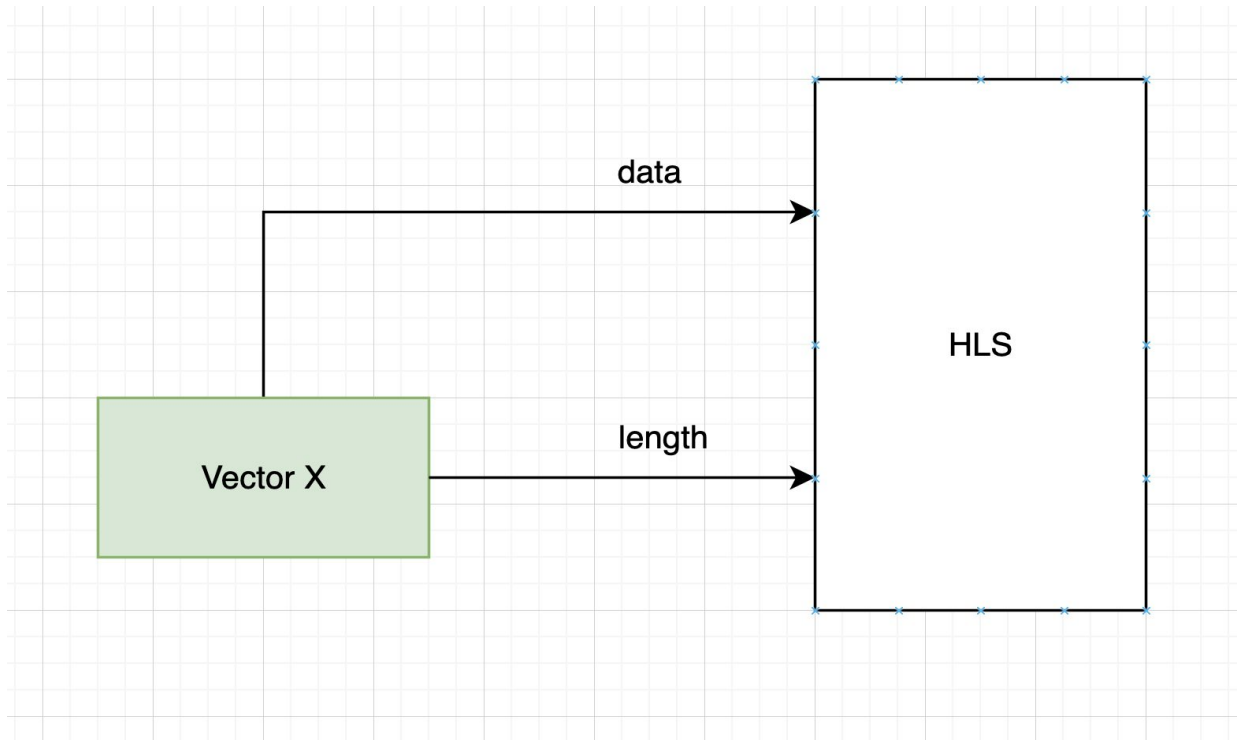
No possible algorithmic optimization!



## Algorithm in C

```
4
5  i32 l2_norm(const i32 a[N]) {
6      i32 tmp, tmp_sum, sum = 0;
7      for (auto i = 0; i < N; i++) {
8          tmp = a[i]*a[i];
9          tmp_sum = tmp + sum;
10         if (tmp_sum < sum) {
11             // Overflow happen, return maximum value then exit
12             sum = INT32_MAX;
13             break;
14         }
15         else {
16             sum = tmp_sum;
17         }
18     }
19     return sum;
20 }
```

# Input Flow





## Reference C

```
4
5  i32 l2_norm(const i32 a[N]) {
6      i32 tmp, tmp_sum, sum = 0;
7      for (auto i = 0; i < N; i++) {
8          tmp = a[i]*a[i];
9          tmp_sum = tmp + sum;
10         if (tmp_sum < sum) {
11             // Overflow happen, return maximum value then exit
12             sum = INT32_MAX;
13             break;
14         }
15         else {
16             sum = tmp_sum;
17         }
18     }
19     return sum;
20 }
```



# Setup

## 1. Setup Testbench

Return 1 if ERROR

Return 0 if CORRECT

```
32
33     e_gold = l2_norm(a);
34     e_hls = hls_l2_norm(a);
35
36     cout << "gold: " << e_gold << endl;
37     cout << " HLS: " << e_hls << endl;
38
39     if (e_gold != e_hls)
40     {
41         cout << "Error" << endl;
42         return 1;
43     }
44
45     cout << "OK" << endl;
46
47     return 0;
48 }
```





# Setup

## 2. Setup Header

Header is used to link the testbench and the HLS code.

```
12
13     typedef ap_int<16> i16;
14     typedef ap_int<32> i32;
15
16
17     i32 hls_l2_norm(const i32 a[N]);
18
19
```



# Setup

## 3. HLS function

HLS function is  
same as C reference code

```
2
3  i32 hls_l2_norm(const i32 a[N])
4  {
5      i32 tmp, tmp_sum, sum = 0;
6      for (auto i = 0; i < N; i++)
7      {
8          tmp = a[i] * a[i];
9          tmp_sum = tmp + sum;
10         if (tmp_sum < sum)
11         {
12             // Overflow happen, return maximum value then exit
13             sum = INT32_MAX;
14             break;
15         }
16         else
17         {
18             sum = tmp_sum;
19         }
20     }
21     return sum;
22 }
```



# Baseline

The baseline is from the C reference implementation

1. Does not have Hardware Interface
2. Is not accelerated
3. Is a start point

**Task: Improve performance of Function in High Level Synthesis**

# Optimization Strategy

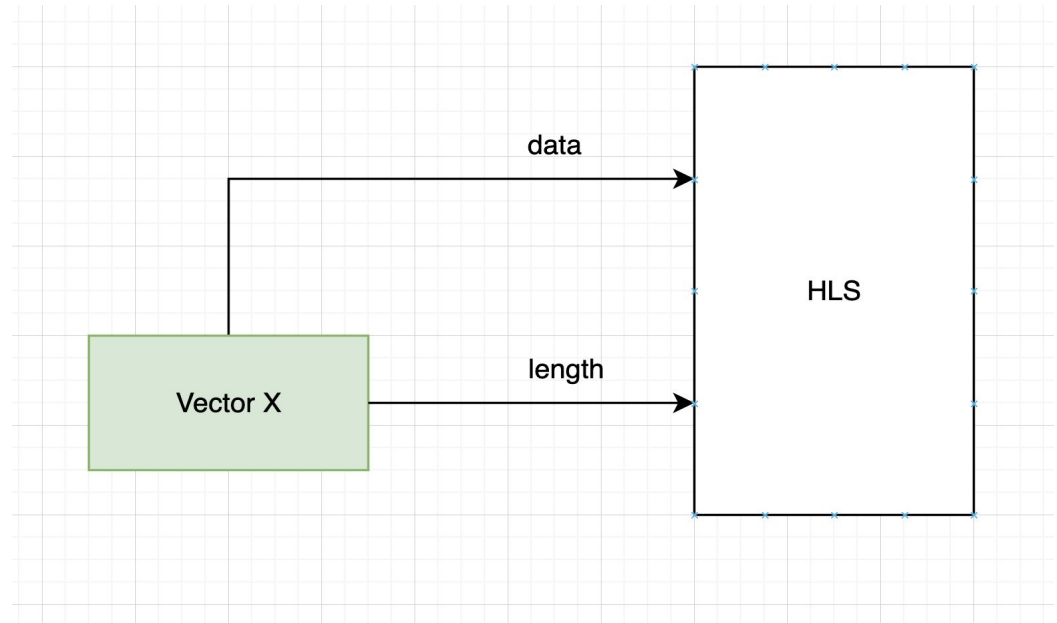
Constrains:

(1) No memory storage

Goal:

(2) Computer as fast as possible

Demo: Read C code



# Optimization Strategy

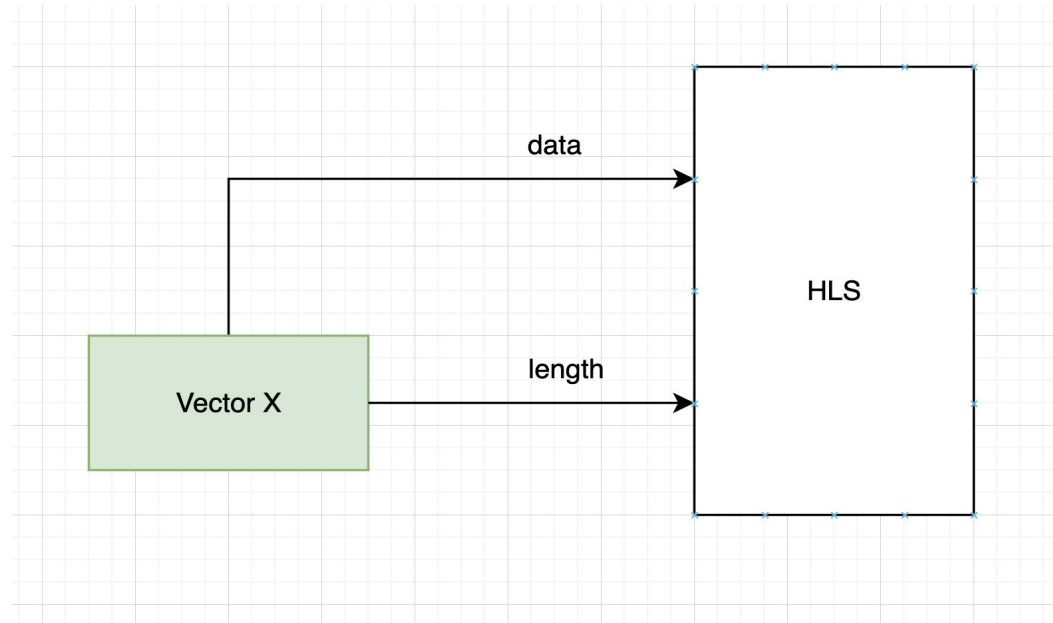
Constrains:

(1) No memory storage

Goal:

(2) Computer as fast as possible

=> We can only increase bus-width





## **Increase from 32-bit to 512-bit**

- Use maximum bus-width available
- In theoretical, if we can proceed 512-bit, we will improve performance 16 times faster



## Increase from 32-bit to 512-bit

- Use maximum bus-width available
- In theoretical, if we can proceed 512-bit, we will improve performance 16 times faster

Demo



## Increase from 32-bit to 512-bit

In summary, we can use two syntaxes to read data:

### Range Selection

```
ap_range_ref ap_(u)int::range (unsigned Hi, unsigned Lo)  
ap_range_ref ap_(u)int::operator () (unsigned Hi, unsigned Lo)
```





## Increase from 32-bit to 512-bit

We can select bit range by () or range():

```
ap_uint<4> Rslt;  
  
ap_uint<8> Val1 = 0x5f;  
ap_uint<8> Val2 = 0xaa;  
  
Rslt = Val1.range(3, 0); // Yields: 0xF  
Val1(3,0) = Val2(3, 0); // Yields: 0x5A  
Val1(3,0) = Val2(4, 1); // Yields: 0x55  
Rslt = Val1.range(4, 7); // Yields: 0xA; bit-reversed!
```



## Increase from 32-bit to 512-bit

- Handle bit-range selection is sometimes confusing
- We can use two dimensional array to represent wider bus-width with `#pragma ARRAY_RESHAPE`



## Increase from 32-bit to 512-bit

- Handle bit-range selection is sometimes confusing
- We can use two dimensional array to represent wider bus-width with `#pragma ARRAY_RESHAPE`



# Multi-dimensional Array wider bus-width

Demo

# #pragma HLS ARRAY\_RESHAPE


factor=2

Left side:

- Spent  $N$  clock cycles to read  $N$  elements in array

Right side:

- **Block:** 1 clock cycle read 2 elements **distance by  $N/2$**



# #pragma HLS ARRAY\_RESHAPE


factor=2

Left side:

- Spent N clock cycles to read N elements in array

Right side:

- **Cyclic:** 1 clock cycle read 2 elements **distance by 1**



# #pragma HLS ARRAY\_RESHAPE


factor=2

Left side:

- Spent N clock cycles to read N elements in array

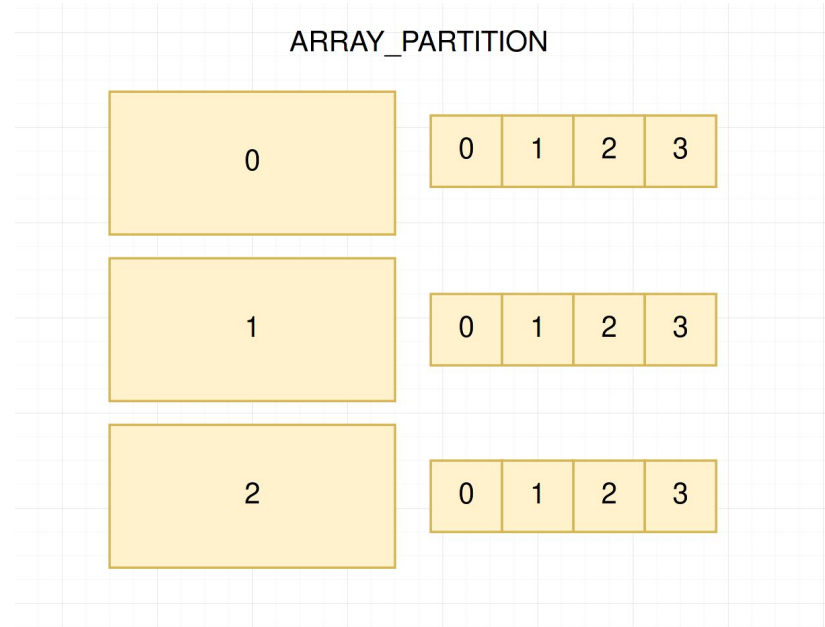
Right side:

- **Complete:** 1 clock cycle read all elements



## #pragma HLS ARRAY\_PARTITION

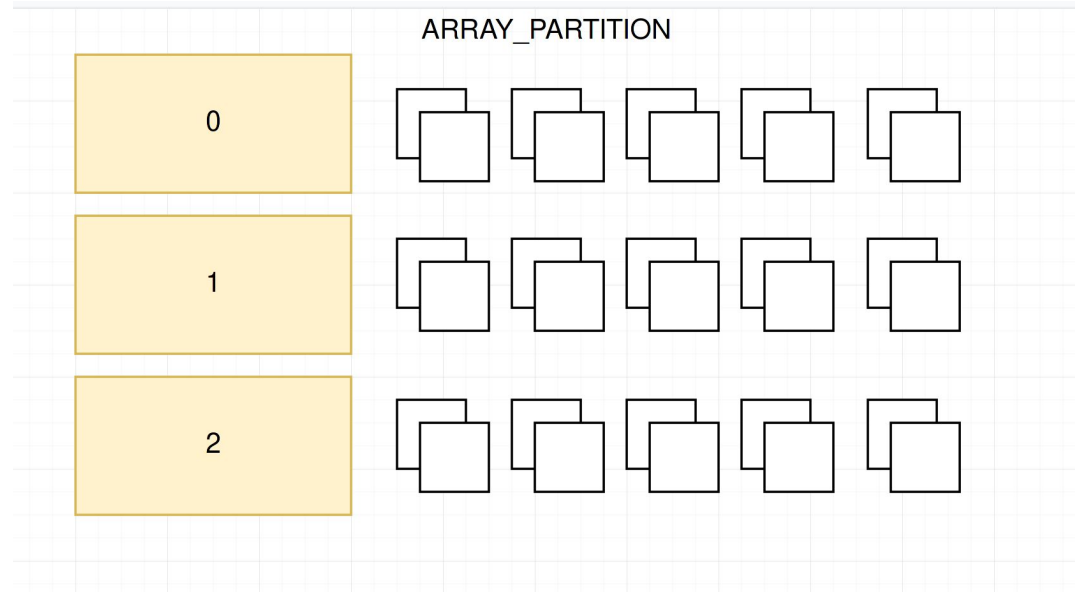
- Similar to  
ARRAY\_RESHAPE
- Easier to control  
multi-dimensional array





## #pragma HLS ARRAY\_PARTITION

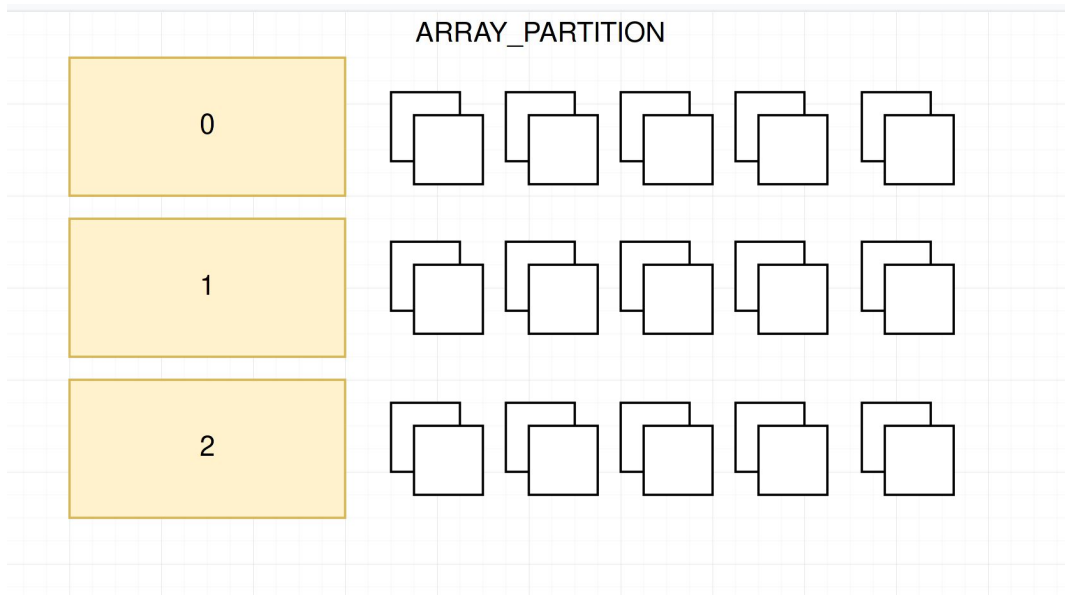
- Similar to  
ARRAY\_RESHAPE
- Easier to control  
multi-dimensional array



## ARRAY\_PARTITION vs ARRAY\_RESHAPE

Two pragma are very similar,  
when to use what ?

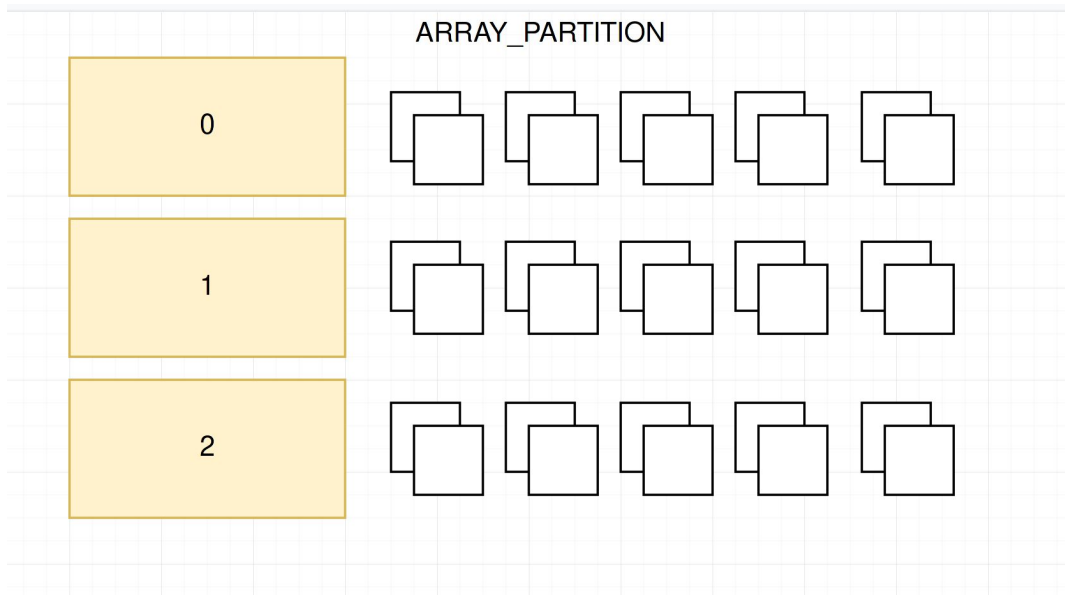
- ARRAY\_RESHAPE for **single** dimensional array



## ARRAY\_PARTITION vs ARRAY\_RESHAPE


Two pragma are very similar,  
when to use what ?

- ARRAY\_RESHAPE for **single** dimensional array
- ARRAY\_PARTITION for **multi-dimensional** array




# Adder Tree

What we have




## Fan-out Adder Tree



Huge adder tree decrease  
performance

# Fanout Problem


- Long critical path






## Fanout Problem

- Long critical path



# Fanout Problem

- Long critical path








# Fanout Problem


- Long critical path



# Adder Tree

What we want:

Multiple layer of  
adder tree




# Partial UNROLL or “Folding”

Simple explanation:

1. Chop big loop to smaller loop

FOLDING



# Partial UNROLL or “Folding”


Simple explanation:

1. Chop big loop to smaller loop

eh...?

2. Chop Chop Chop

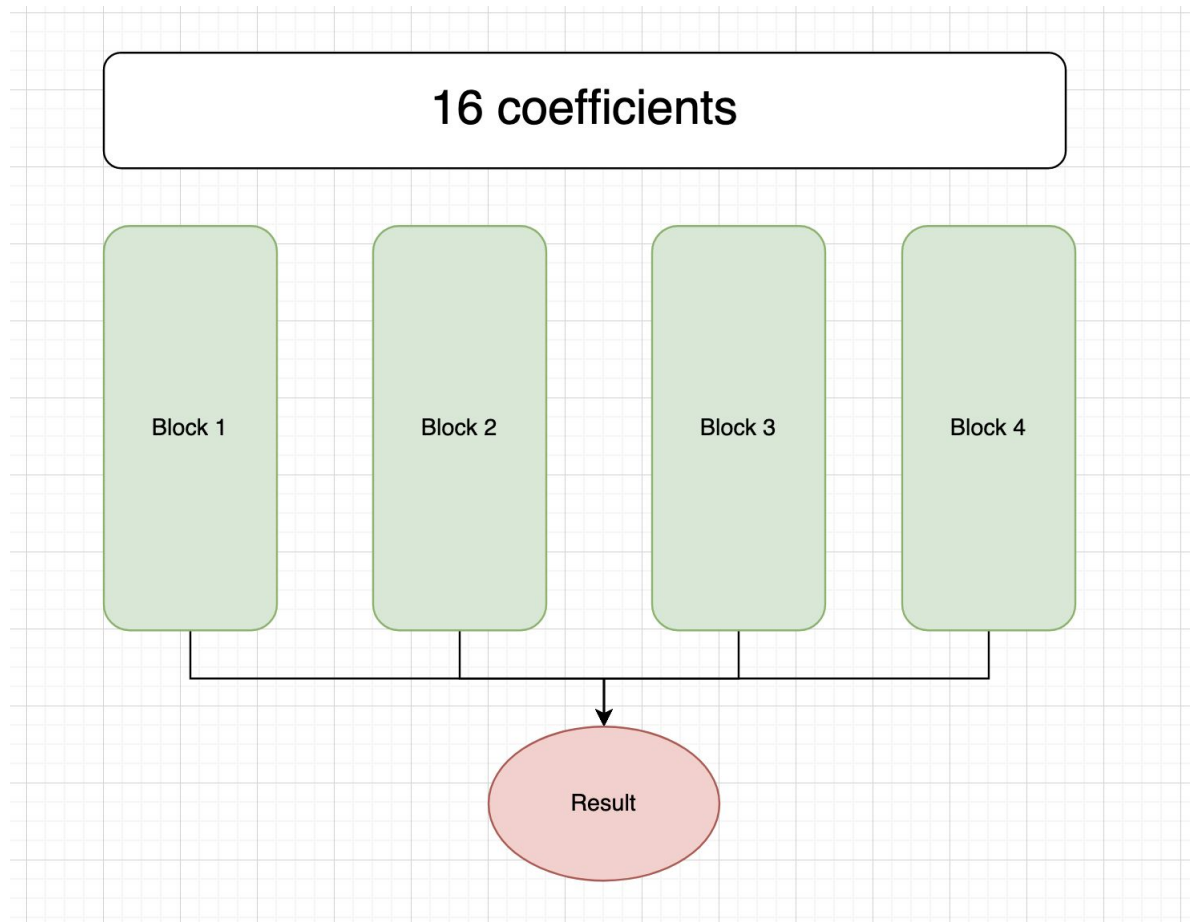
FOLDING



# Folding Demo

Demo

-





## Decrease loop performance

```
11 // First stage adder
12 adder_tree_1:
13     for (auto i = 0; i < 4; i++)
14     {
15         tmp[i] = 0;
16         for (auto j = 0; j < 4; j++)
17         {
18             tmp[i] += a[i * 4 + j] * a[i * 4 + j];
19         }
20     }
21
```

What the heck is this?

## Decrease loop performance

```
11 // First stage adder
12 adder_tree_1:
13     for (auto i = 0; i < 4; i++)
14     {
15         tmp[i] = 0;
16         for (auto j = 0; j < 4; j++)
17         {
18             tmp[i] += a[i * 4 + j] * a[i * 4 + j];
19         }
20     }
21
```

What the heck is this?




# #pragma HLS LOOP\_FLATTEN

Either:

- In every iteration of the **outer loop**, it spends 1 cycles to jump into **inner loop**

Or

- After **inner loop** finished, it spends 1 cycles to jump to **outer loop**






# #pragma HLS LOOP\_FLATTEN

Either:


- In every iteration of the **outer loop**, it spends 1 cycles to jump into **inner loop**

Or

- After **inner loop** finished, it spends 1 cycles to jump to **outer loop**



# #pragma HLS LOOP\_FLATTEN





## Should we assign value outside the loop?

- If you assign value at the outer loop, you have to spend **1 cycles** to get to the inner loop.
- If you assign value at the initialization, you rely on **HLS compiler** to merge the initialization to the loop.

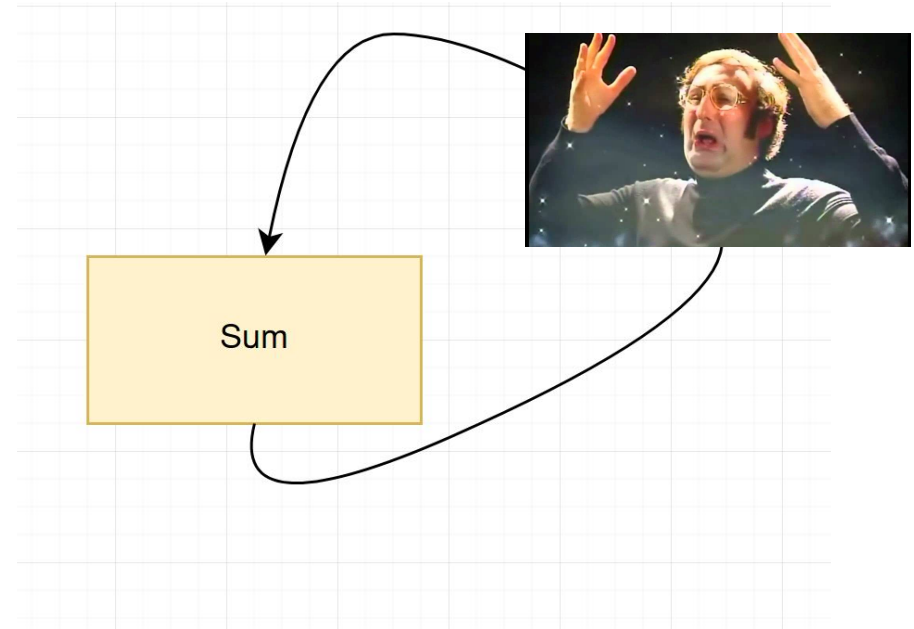
We can do better.

Demo

# Violation

Resolve ADD Violation

Demo






# **We forgot to handle overflow !!!**

Demo

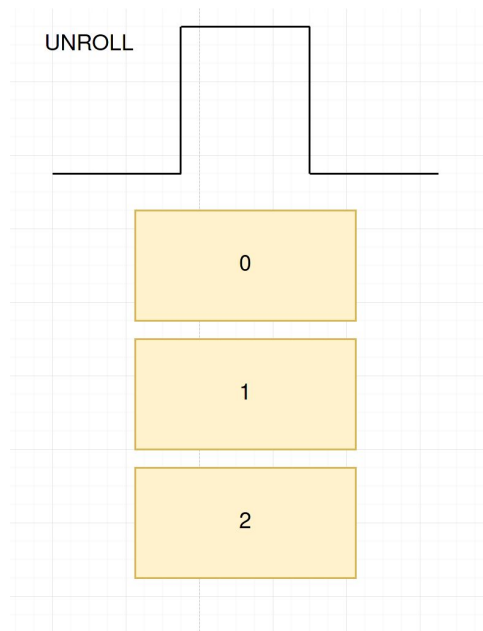
# #pragma HLS UNROLL

Motivation




## #pragma HLS UNROLL

- Parallel execution in **single** clock cycle
- Maximum performance



# Problem with #pragma HLS UNROLL


- Fan-out
- Decrease frequency significantly
- Resources blow up






# #pragma HLS UNROLL

Demo




## #pragma HLS PIPELINE



## #pragma HLS BIND\_OP


Pin specified operation to premade fabric: DSP, BRAM, URAM, etc...



## #pragma HLS BIND\_OP

Pin specified operation to premade fabric: DSP, BRAM, URAM, etc...

Force Pipeline inside fabric with specified latency (auto is good enough)




## #pragma HLS BIND\_OP

Pin specified operation to premade fabric: DSP, BRAM, URAM, etc...

Force Pipeline inside fabric with specified latency (auto is good enough)

Boost maximum frequency

CLB ? Who cares



# #pragma HLS BIND\_OP

Pin specified operation to premade fabric: DSP, BRAM, URAM, etc...


Force Pipeline inside fabric with specified latency (auto is good enough)

Boost maximum frequency

Demo

- C synthesis

CLB ? Who cares





## Coding Conclusion

- If you are looking at lightweight implementation: choose HDL
  - HLS is an abstract layer over HDL
  - HLS can perform complex operation as HDL
- 
- Pr with HLS is easy
  - Testing with HLS is quick
  - Co-simulation with HLS is quick



## Performance Conclusion

- HLS can achieve good performance.
- CERG GMU demonstrates HLS can achieve similar latency, frequency as RTL.

But HLS utilizes more resources compare to RTL.

- The problem with FPGA in HPC nowadays is not to minimize LUT, FF usage.

It is to achieve **maximum performance**, utilize all available resources on FPGA.





## Performance Conclusion

- HLS can achieve good performance.
- CERG GMU demonstrates HLS can achieve similar latency, frequency as RTL. HLS utilize more resources compare to RTL.
- The problem with FPGA in HPC nowadays is not to minimize LUT, FF usage. It is to achieve maximum performance, utilize all available resources on FPGA.
- Remember to take the transfer size into account.
- **Prepare input data can give ultimate solution!!!**



# What the point of Hardware accelerator anyway?

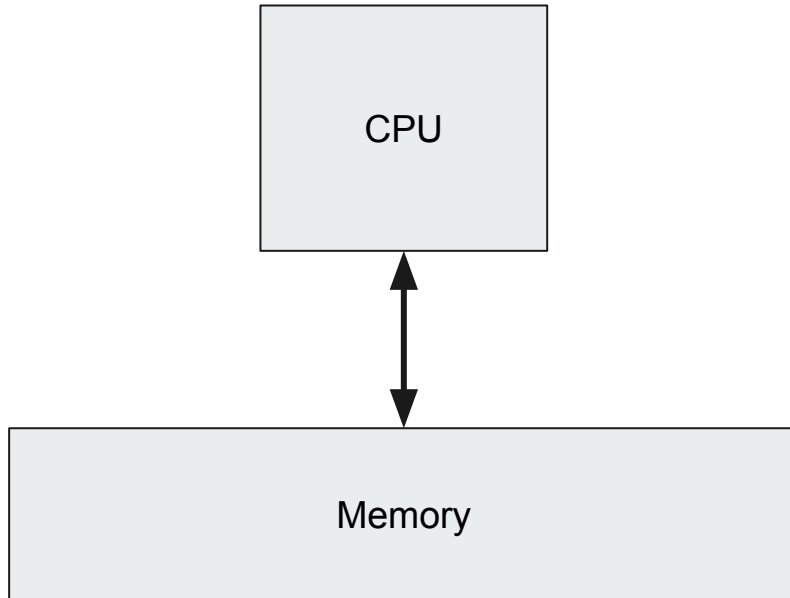
- Not mentioned in this talk: `#pragma HLS DATAFLOW`
- Not mentioned in this talk: NEON vs HLS
- Load/Store architecture and Hardware Accelerator
- The code for NEON is ready. It's up to you to figure it out which one is faster in this case. And let the class know!!!

# NEON Implementation

```

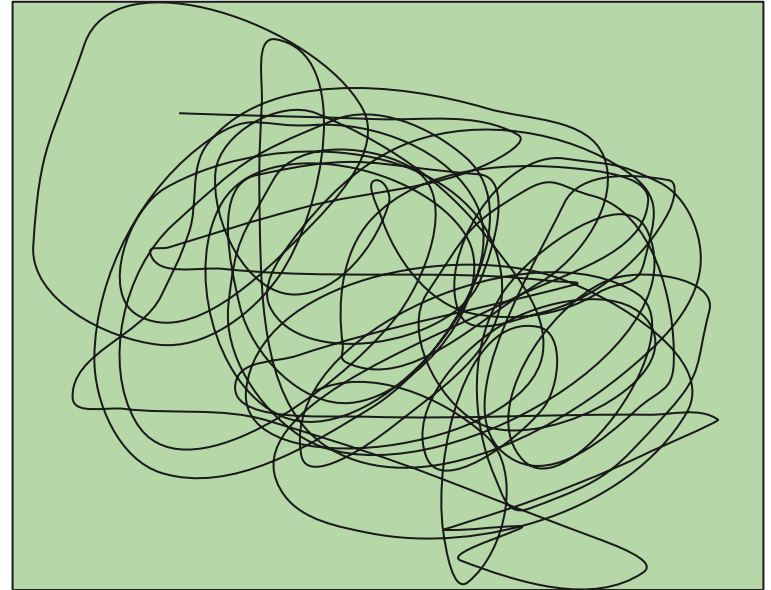
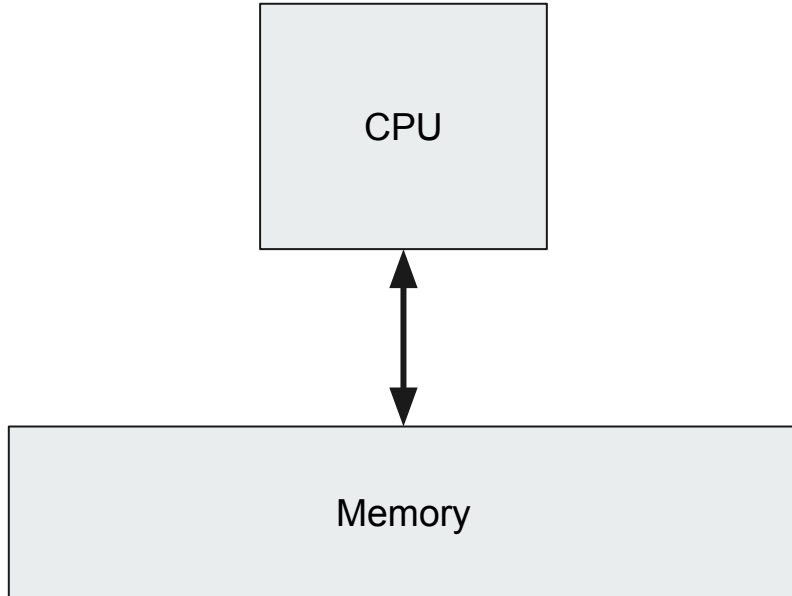
;-- FUNC.100003dc0:
; UNKNOWN XREF from aav.0x10000020 @ +0xb0
92: sym.neon_l2_norm_short_const_ ();
    0x100003dc0    00e4006f    movi v0.2d, 0000000000000000 ; [00] -r-x section size 428 na
med 0.__TEXT.__text
    0x100003dc4    e8038092    movn x8, 0x1f
    0x100003dc8    01e4006f    movi v1.2d, 0000000000000000
    0x100003dcc    0224df4c    ld1 {v2.8h, v3.8h, v4.8h, v5.8h}, [x0], 0x40
    0x100003dd0    4090620e    sqdmlal v0.4s, v2.4h, v2.4h
    0x100003dd4    6090630e    sqdmlal v0.4s, v3.4h, v3.4h
    0x100003dd8    8090640e    sqdmlal v0.4s, v4.4h, v4.4h
    0x100003ddc    a090650e    sqdmlal v0.4s, v5.4h, v5.4h
    0x100003de0    4190624e    sqdmlal2 v1.4s, v2.8h, v2.8h
    0x100003de4    6190634e    sqdmlal2 v1.4s, v3.8h, v3.8h
    0x100003de8    8190644e    sqdmlal2 v1.4s, v4.8h, v4.8h
    0x100003dec    a190654e    sqdmlal2 v1.4s, v5.8h, v5.8h
    0x100003df0    08810091    add x8, x8, 0x20
    0x100003df4    1f8107f1    cmp x8, 0x1e0
    0x100003df8    a3fef54    b.lo 0x100003dcc
    0x100003dfc    0004a14e    shadd v0.4s, v0.4s, v1.4s
    0x100003e00    0140006e    ext v1.16b, v0.16b, v0.16b, 8
    0x100003e04    000ca10e    sqadd v0.2s, v0.2s, v1.2s
    0x100003e08    083c0c0e    mov w8, v0.s[1]
    0x100003e0c    0101271e    fmov s1, w8
    0x100003e10    000ca15e    sqadd s0, s0, s1
    0x100003e14    0000261e    fmov w0, s0
    0x100003e18    c0035fd6    ret
[0x100003dc0]>
```

## Load/Store Architecture vs Hardware Accelerator





## Load/Store Architecture vs Hardware Accelerator

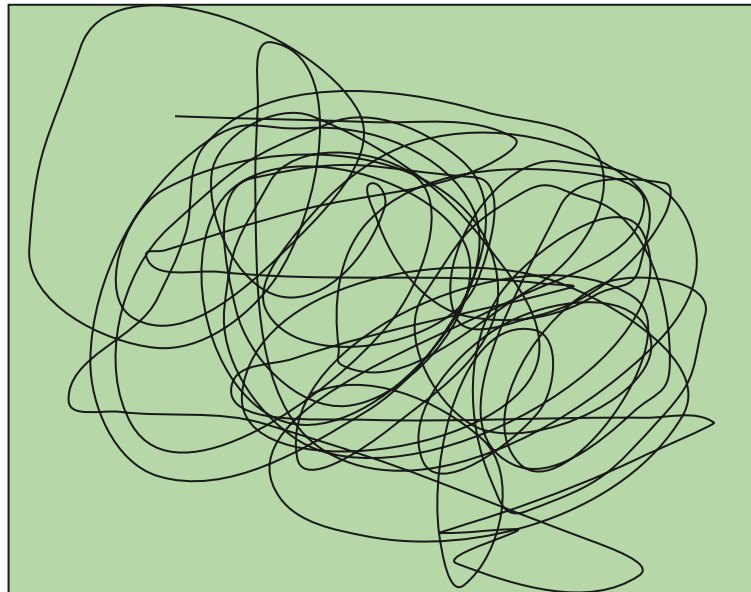


# Load/Store Architecture vs Hardware Accelerator

- No pressure on CPU
- No Store and Load

Real world example:

- Video decoder: There are plenty optimized SW implementations. If you don't enable HW, it's give low framerate and high CPU usage.





# Thank you

Question ???



# Reference

Code used in this talk:

- [https://github.com/cothan/Vitis\\_High\\_Level\\_Synthesis\\_Training](https://github.com/cothan/Vitis_High_Level_Synthesis_Training)
- [https://github.com/cothan/Vitis\\_High\\_Level\\_Synthesis\\_Training\\_Part\\_2](https://github.com/cothan/Vitis_High_Level_Synthesis_Training_Part_2)
- 

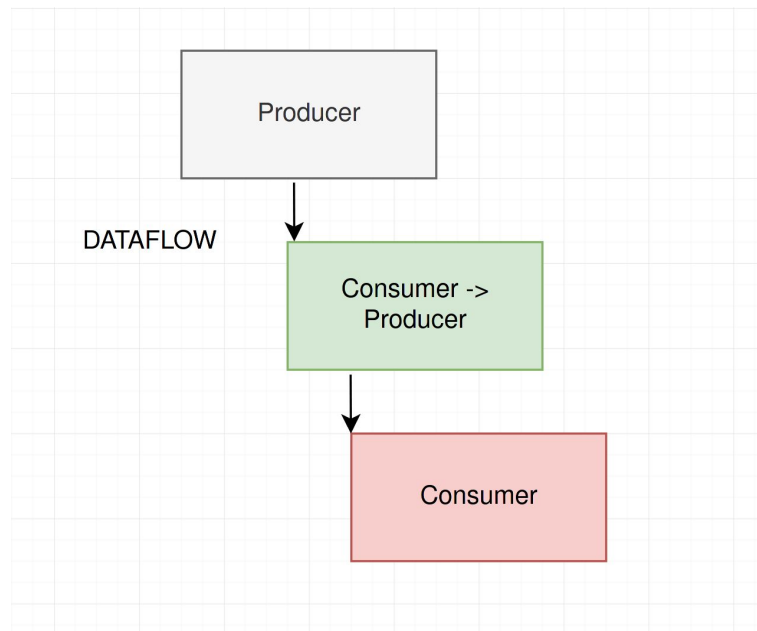
Helpful documents which are not from Xilinx:

- Transformations of High-Level Synthesis Codes for High-Performance Computing  
<https://arxiv.org/abs/1805.08288>



# #pragma HLS DATAFLOW

DATAFLOW need to be written in  
producer-consumer fashion



# #pragma HLS DATAFLOW

DATAFLOW need to be written in producer-consumer fashion

May increase or decrease the resources utilization.

